



# CS 6220 Data Mining — Assignment 1

## Due: September 19, 2023(100 points)

---

YOUR NAME + LDAP  
[link-to-your-repository](#)

## Review of Prerequisite Coding

In subsequent lectures, you'll learn about frequent item sets, where relationships between items are learned by observing how often they co-occur in a set of data. This information is useful for making recommendations in a rule based manner. Before looking at frequent item sets, it is worth understanding the space of all possible sets and get a sense for how quickly the number of sets with unique items grows.

Suppose that we've received only a hundred records of items bought by customers at a market. Each line in the file represents the items an individual customer bought, i.e. their basket. For example, consider the following rows.

```
ham, cheese, bread
dates, bananas
celery, chocolate bars
```

Customer 1 has a basket of ham, cheese, and bread. Customer 2 has a basket of dates and bananas. Customer 3 has a basket of celery and chocolate bars. Each of these records is the receipt of a given customer, identifying what they bought.

**Please answer the following:**

1. The *cardinality* of a set or collection of items is the number of unique items in that collection. Write a function called `cardinality_items` that takes a `.csv` text string file as input, where the format is as the above, and calculates the cardinality of the set of all the grocery items in any given dataset. What is the cardinality in “`basket.data.csv`”?
2. We'd occasionally like to understand the space of all possible subsets comprised of unique items. If there are  $N$  unique items (i.e., the cardinality of the entire dataset is  $N$ ), how

many sets with unique items can there possibly be? (Ignore the null set.) NOTE: I only expect the formula, and there is no code associated with this question.

3. Write a function called `all_itemsets` that takes a list of unique items and an integer  $N$  as input, and the output is a list of all possible unique itemsets with non-repeating  $N$  items. That is, the output is  $L = [S_1, S_2, \dots, S_M]$ , a list of all possible sets of  $N$  unique items.

For example,

```
all_itemsets( ["ham", "cheese", "bread"], 2 )
```

should result in:

```
[ ["ham", "cheese"], ["ham", "bread"], ["cheese", "bread"] ]
```

You should not need any library functions.

4. Let's take the “[basket.data.csv](#)” file and try to determine the probability of a given set occurring. For example, if the set  $S = \{\text{bread}, \text{oatmeal}\}$  occurs twice in a dataset with 100 records, then the probability of item set  $\{\text{bread}, \text{oatmeal}\}$  occurring is 0.02. Likewise in another example, if  $S = [\text{"dates"}, \text{"bananas"}]$  and the data file looks like :

```
ham, cheese, bread
dates, bananas
celery, chocolate bars
```

then `prob_S("data_file.csv", S)` should yield 0.3333 since this set occurs a third of the time in the data file.

Write a function called `prob_S` that takes a `.csv` text string file as input and an item set (e.g.,  $S = \{\text{"bread"}, \text{"oatmeal"}\}$ ) and outputs the probability of seeing  $S$ .

## Examining Our First Dataset

One of the most famous challenges in data science and machine learning is Netflix's Grand Prize Challenge, where Netflix held an open competition for the best algorithm to predict user ratings for films. The grand prize was \$1,000,000 and was won by BellKor's Pragmatic Chaos team. This is the dataset that was used in that competition.

- <https://www.kaggle.com/datasets/netflix-inc/netflix-prize-data>

In this exercise, we're going to do a bit of exploring in the Netflix Data. Start by downloading the data. Data integrity tends to be a problem in large scale processing, especially if there is little to no support. Therefore, it's important to verify the quality of the file download. If all worked out well, you should have the following files:

- combined\_data\_1.txt
- combined\_data\_2.txt
- combined\_data\_3.txt
- combined\_data\_4.txt
- movie\_titles.csv
- probe.txt
- qualifying.txt
- README

## Data Verification and Analysis

A large part of machine learning and data science is about getting data in the right format and ensuring that there is no data corruption. Verify that the schema is the same as the Kaggle Dataset's description, read in the data (hint: some movies have commas in them), and then answer the following questions.

**Please answer the following:**

5. How many total records of movie ratings are there in the entire dataset (over all of combined\_data\_\*.txt)? How many total users are there in the entire dataset (over all of combined\_data\_\*.txt)?
6. How many movies with unique names are there?
7. What is the range of years that this data is valid over?
8. Year over year, what is the trend of the number of ratings? Plot the number of movie ratings per year (aggregated over all users). Any guesses as to why this is the case?
9. How many users rated exactly 200 movies? Of these users, take the lowest user ID and print out the names of the movies that they liked the most (5 star ratings).